

Boston House Price Prediction

Akhil

2023-24-4

Loading Packages

```
library(mlbench)
library(ggplot2)
library(beeswarm)
```

dataset description

crim | per capita crime rate by town
zn | proportion of residential land zoned for lots over 25,000 sq.ft
indus | proportion of non-retail business acres per town
chas | Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
nox | nitric oxides concentration (parts per 10 million)
rm | average number of rooms per dwelling
age | proportion of owner-occupied units built prior to 1940
dis | weighted distances to five Boston employment centres
rad | index of accessibility to radial highways
tax | full-value property-tax rate per USD 10,000
ptratio | pupil-teacher ratio by town
b | where BB is the proportion of blacks by town
lstat | percentage of lower status of the population
medv | median value of owner-occupied homes in USD 1000's | target variable

Data Input

```
data(BostonHousing)
head(BostonHousing)
```

```
##      crim zn  indus chas   nox    rm  age    dis rad tax ptratio    b lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296   15.3 396.90  4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242   17.8 396.90  9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242   17.8 392.83  4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222   18.7 394.63  2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222   18.7 396.90  5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222   18.7 394.12  5.21
##   medv
## 1 24.0
## 2 21.6
```

```
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7
```

```
str(BostonHousing)
```

```
## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : num 1 2 2 3 3 3 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ b : num 397 397 393 395 397 ...
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
# checking for null values
colSums(is.na(BostonHousing))
```

```
##      crim      zn      indus      chas      nox      rm      age      dis      rad      tax
##      0        0        0        0        0        0        0        0        0        0
## ptratio      b      lstat      medv
##      0        0        0        0
```

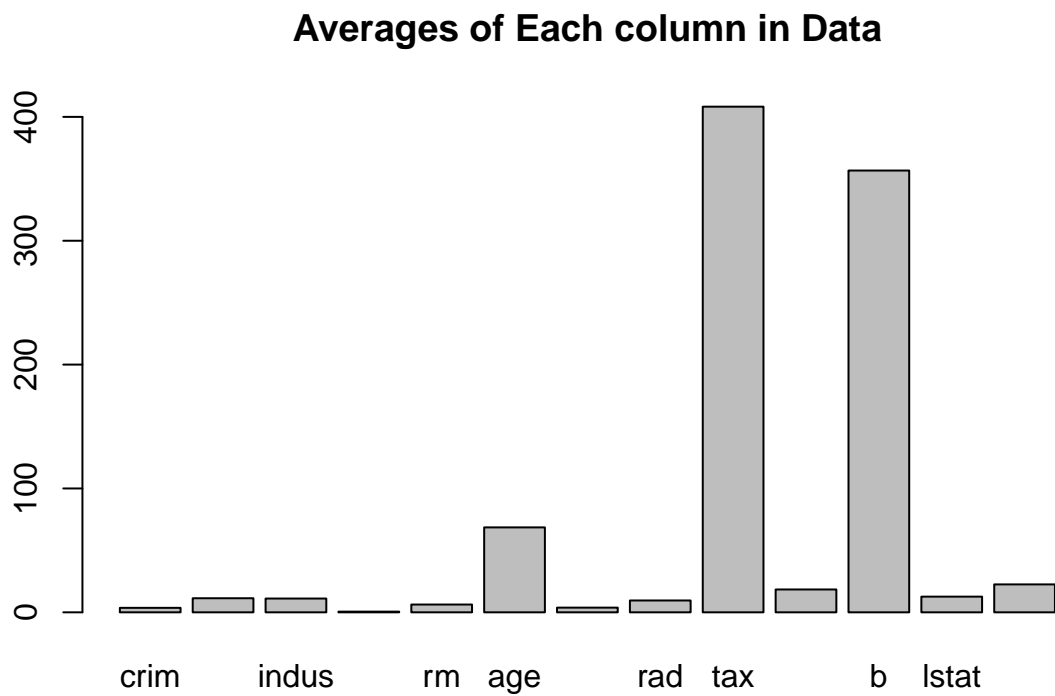
```
summary(BostonHousing)
```

```
##      crim      zn      indus      chas      nox
## Min.   : 0.00632 Min.   : 0.00 Min.   : 0.46 0:471 Min.   :0.3850
## 1st Qu.: 0.08205 1st Qu.: 0.00 1st Qu.: 5.19 1: 35 1st Qu.:0.4490
## Median : 0.25651 Median : 0.00 Median : 9.69 Median :0.5380
## Mean   : 3.61352 Mean   : 11.36 Mean   :11.14 Mean   :0.5547
## 3rd Qu.: 3.67708 3rd Qu.: 12.50 3rd Qu.:18.10 3rd Qu.:0.6240
## Max.   :88.97620 Max.   :100.00 Max.   :27.74 Max.   :0.8710
##      rm      age      dis      rad
## Min.   :3.561 Min.   : 2.90 Min.   : 1.130 Min.   : 1.000
## 1st Qu.:5.886 1st Qu.: 45.02 1st Qu.: 2.100 1st Qu.: 4.000
## Median :6.208 Median : 77.50 Median : 3.207 Median : 5.000
## Mean   :6.285 Mean   : 68.57 Mean   : 3.795 Mean   : 9.549
## 3rd Qu.:6.623 3rd Qu.: 94.08 3rd Qu.: 5.188 3rd Qu.:24.000
## Max.   :8.780 Max.   :100.00 Max.   :12.127 Max.   :24.000
##      tax      ptratio      b      lstat
## Min.   :187.0 Min.   :12.60 Min.   : 0.32 Min.   : 1.73
## 1st Qu.:279.0 1st Qu.:17.40 1st Qu.:375.38 1st Qu.: 6.95
## Median :330.0 Median :19.05 Median :391.44 Median :11.36
## Mean   :408.2 Mean   :18.46 Mean   :356.67 Mean   :12.65
```


```
## 3rd Qu.:666.0 3rd Qu.:20.20 3rd Qu.:396.23 3rd Qu.:16.95
## Max. :711.0 Max. :22.00 Max. :396.90 Max. :37.97
## medv
## Min. : 5.00
## 1st Qu.:17.02
## Median :21.20
## Mean :22.53
## 3rd Qu.:25.00
## Max. :50.00
```

Visualization

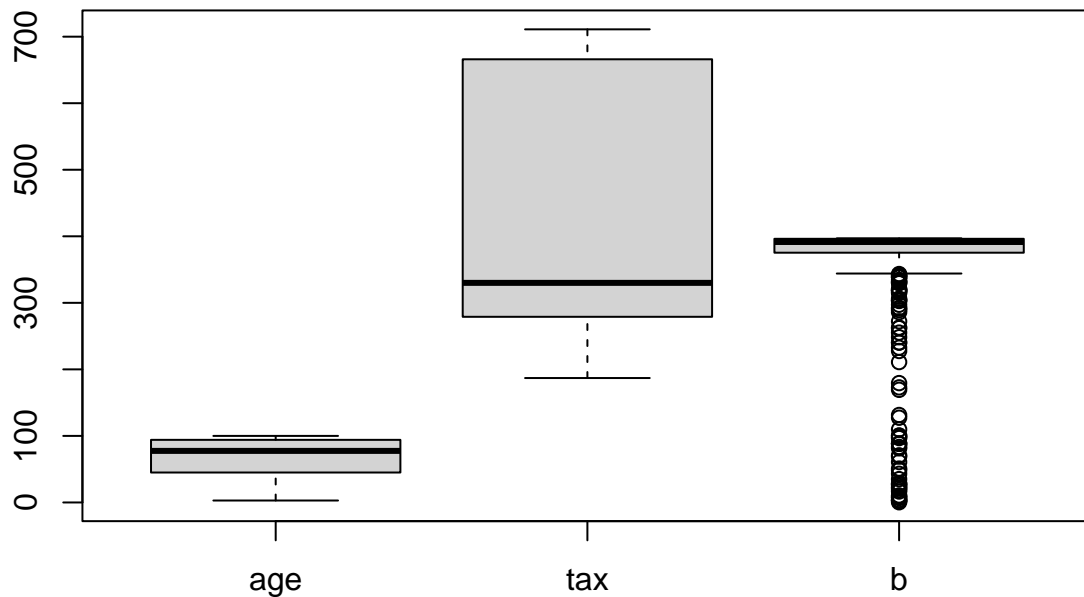
```
avg<-colMeans(BostonHousing[,-4])
barplot(avg, main="Averages of Each column in Data")
```



```
boxplot(BostonHousing[-c(4,7,10,12)])
```



```
boxplot(BostonHousing[c(7,10,12)])
```

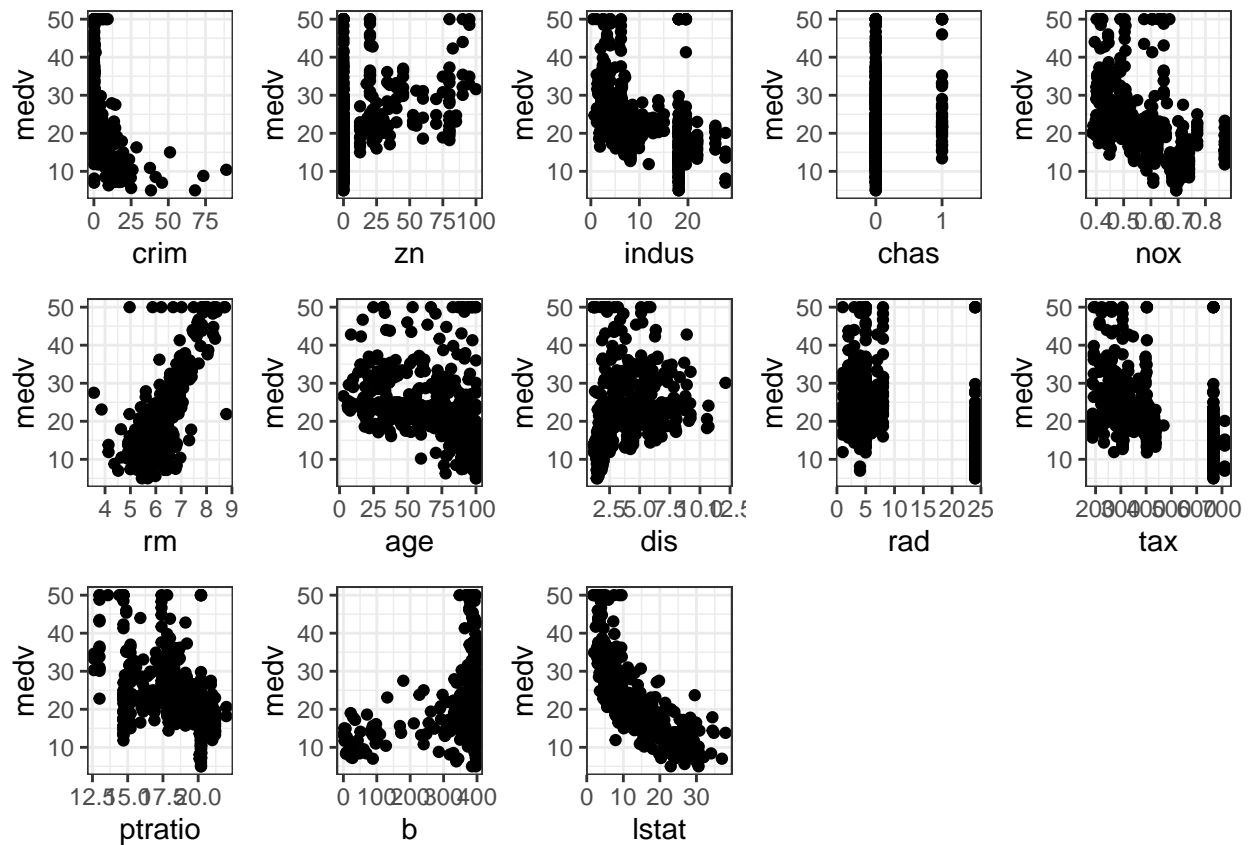


```
# choosing to plot medv vs remaining columns in dataset
price <- "medv"

# create a for loop to iterate over each column in the data frame
plots_list <- lapply(names(BostonHousing), function(var) {
  if (var != price) {
    ggplot(BostonHousing, aes_string(x = var, y = price)) +
      geom_point() +
      labs(x = var, y = price) +
      theme_bw()
  } else {
    NULL
  }
})
```


```
## Warning: 'aes_string()' was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation ideoms with 'aes()'
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

```
library(gridExtra)
grid.arrange(grobs = plots_list, ncol = 5)
```




Plot between Average Room vs Price


```
ggplot(BostonHousing) +
  aes(x = rm, y = medv) +
  geom_point(shape = "circle", size = 1.5, colour = "#112446") +
  theme_minimal()
```



```
ggplot(BostonHousing) +  
  aes(x = medv, y = tax) +  
  geom_point(shape = "circle", size = 1.5, colour = "#112446") +  
  labs(x = "Price", y = "Tax", title = "Price vs Tax") +  
  theme_minimal()
```



```
beeswarm(BostonHousing$medv~BostonHousing$chas)
```





```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
# correlation plot using the corrplot function
```

```
corrplot(cor(BostonHousing[-4]), method = "color", type = "upper", order = "hclust", addCoef.col = "black")
```



Model building

```
model_bh <- lm(medv~.,data=BostonHousing)
anova(model_bh)
```

```
## Analysis of Variance Table
##
## Response: medv
##          Df Sum Sq Mean Sq  F value    Pr(>F)
## crim      1  6440.8   6440.8  286.0300 < 2.2e-16 ***
## zn        1  3554.3   3554.3  157.8452 < 2.2e-16 ***
## indus     1  2551.2   2551.2  113.2984 < 2.2e-16 ***
## chas      1   1529.8   1529.8   67.9393 1.543e-15 ***
## nox       1    76.2    76.2    3.3861 0.0663505 .
## rm        1 10938.1  10938.1  485.7530 < 2.2e-16 ***
## age       1    90.3    90.3    4.0087 0.0458137 *
## dis       1  1779.5   1779.5   79.0262 < 2.2e-16 ***
## rad       1    34.1    34.1    1.5159 0.2188325
## tax       1   329.6    329.6   14.6352 0.0001472 ***
## ptratio   1  1309.3   1309.3   58.1454 1.266e-13 ***
## b         1   593.3    593.3   26.3496 4.109e-07 ***
## lstat     1  2410.8   2410.8  107.0634 < 2.2e-16 ***
## Residuals 492 11078.8    22.5
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Selecting variables based on significance

```
library(ISLR)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following object is masked from 'package:gridExtra':
##
##      combine

## The following objects are masked from 'package:stats':
##
##      filter, lag

## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union

boston_filtered <- BostonHousing %>% select("medv", "crim", "zn", "indus", "chas", "rm",
                                           "age", "dis", "tax", "ptratio", "b",
                                           "lstat")
```

Partitioning data into train and test

```
set.seed(123)
library(caret)

## Loading required package: lattice

#Partitioning Data into 80% Training and 20% Validation
Index_Train<-createDataPartition(boston_filtered$medv, p=0.8, list=FALSE)
boston_Train <-boston_filtered[Index_Train,]
boston_Validation <-boston_filtered[-Index_Train,]
```

Normalizing data

```
norm_model<-preProcess(boston_Train, method = c("center", "scale"))
#Applying Normalization model to all three data
boston_norm_Train <-predict(norm_model,boston_Train)
boston_norm_Validation <-predict(norm_model,boston_Validation)
```

Linear Regression Model

```
linear <- lm(medv~.,data=boston_norm_Train)
summary(linear)

##
## Call:
## lm(formula = medv ~ ., data = boston_norm_Train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.78665 -0.29548 -0.07201  0.17752  2.97981
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.021381   0.027964  -0.765   0.44497
## crim        -0.043119   0.035289  -1.222   0.22248
## zn           0.090186   0.040694   2.216   0.02725 *
## indus       -0.106026   0.050236  -2.111   0.03544 *
## chas1        0.290071   0.105482   2.750   0.00623 **
## rm           0.308543   0.036711   8.405 7.86e-16 ***
## age         -0.024971   0.046554  -0.536   0.59199
## dis         -0.270850   0.052204  -5.188 3.40e-07 ***
## tax         -0.007578   0.047761  -0.159   0.87401
## ptratio     -0.162151   0.034229  -4.737 3.03e-06 ***
## b            0.100759   0.031249   3.224  0.00137 **
## lstat       -0.441561   0.047523  -9.292 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5419 on 395 degrees of freedom
## Multiple R-squared:  0.7143, Adjusted R-squared:  0.7063
## F-statistic: 89.78 on 11 and 395 DF,  p-value: < 2.2e-16
```

Decision Tree

```
library(rpart.plot)

## Loading required package: rpart

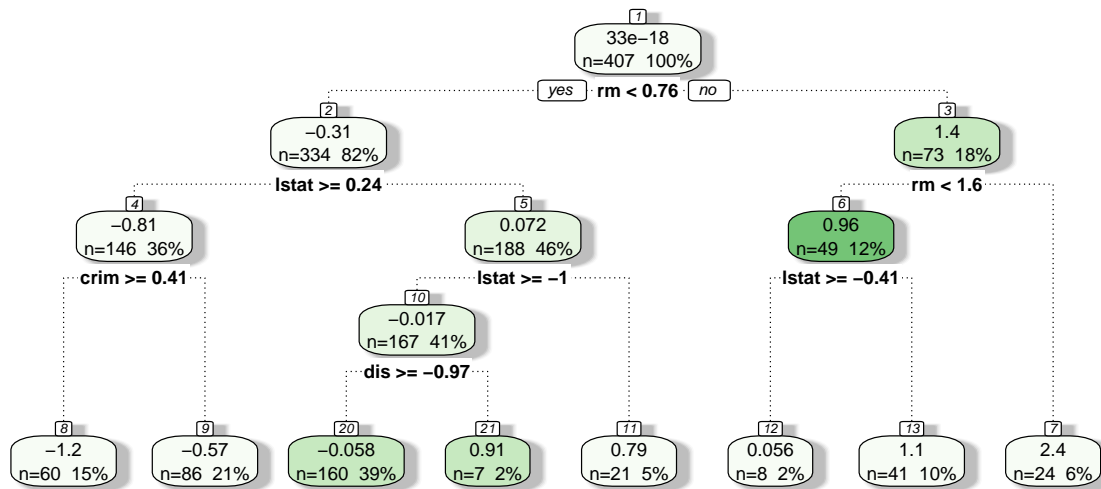
library(rattle)

## Loading required package: tibble

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(rpart)
DT=rpart(medv~.,data=boston_norm_Train, method='anova')
fancyRpartPlot(DT)
```



Rattle 2023-May-07 23:05:18 daraa

```
DT_train <- caret::train(medv~.,data=boston_norm_Train,
                          method = "rpart" )
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

DT_train

```
## CART
##
## 407 samples
## 11 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 407, 407, 407, 407, 407, 407, ...
## Resampling results across tuning parameters:
##
##   cp          RMSE      Rsquared   MAE
## 0.08534102 0.6604279 0.5682854 0.4723816
```

```
## 0.15865427 0.7284011 0.4721545 0.5323949
## 0.45374690 0.8093380 0.3802377 0.6012841
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was cp = 0.08534102.
```

Random Forest

```
set.seed(123)
Random_forest<-train(medv~., data=boston_norm_Train,method='rf')
print(Random_forest)
```

```
## Random Forest
##
## 407 samples
## 11 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 407, 407, 407, 407, 407, 407, ...
## Resampling results across tuning parameters:
##
## mtry RMSE Rsquared MAE
## 2 0.4177788 0.8369361 0.2781526
## 6 0.3754063 0.8600483 0.2590257
## 11 0.3921685 0.8445687 0.2707028
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 6.
```

SVM

```
set.seed(123)
svm<-train(medv~., data=boston_norm_Train,method='svmLinear')
print(svm)
```

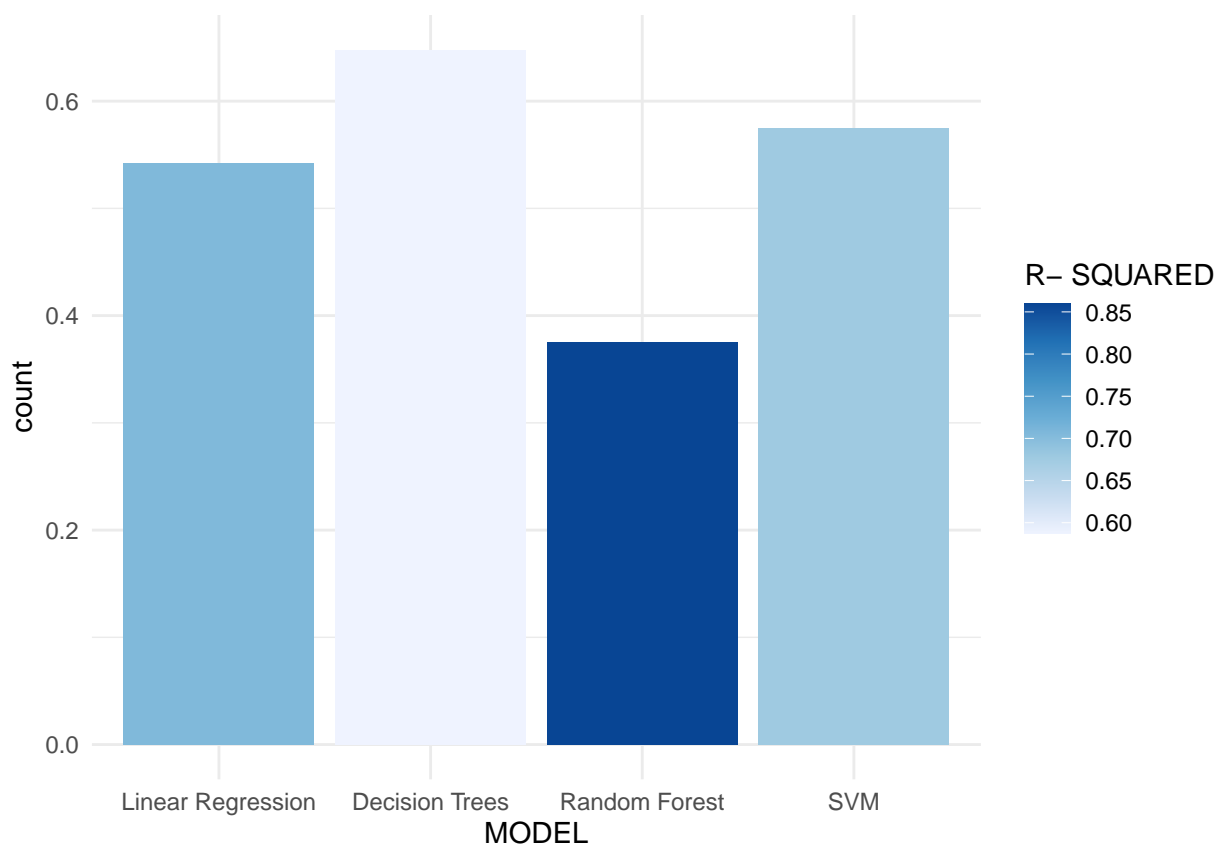
```
## Support Vector Machines with Linear Kernel
##
## 407 samples
## 11 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 407, 407, 407, 407, 407, 407, ...
## Resampling results:
##
## RMSE Rsquared MAE
## 0.5744627 0.6776097 0.3528924
##
## Tuning parameter 'C' was held constant at a value of 1
```

Selecting Best Model

```
library(readr)
results <- read_csv("results.csv", col_types = cols(MODEL = col_factor(levels = c("Linear Regression",

## Warning: One or more parsing issues, call 'problems()' on your data frame for details,
## e.g.:
##   dat <- vroom(...)
##   problems(dat)

results <- na.omit(results)
ggplot(results) + aes(x = MODEL, fill = `R- SQUARED`, weight = RMSE) + geom_bar() +
  scale_fill_distiller(palette = "Blues", direction = 1) + theme_minimal()
```



Based on the R- squared values from the above models, it can be seen that Random forest is performing good with more than 86%

Predicting on test data

```
set.seed(123)
Random_forest_2<-train(medv~., data=boston_Train,method='rf',
  preProcess = c("center", "scale"))
predicted <- c(predict(Random_forest_2,boston_Validation[-1]))
```

```
output<-as.data.frame(predicted)
output$actual <- boston_Validation$medv
```

```
output
```

| ## | predicted | actual |
|--------|-----------|--------|
| ## 3 | 35.106687 | 34.7 |
| ## 6 | 25.146390 | 28.7 |
| ## 9 | 18.358760 | 16.5 |
| ## 11 | 21.102163 | 15.0 |
| ## 14 | 20.209577 | 20.4 |
| ## 15 | 19.938947 | 18.2 |
| ## 31 | 14.604677 | 12.7 |
| ## 32 | 20.652347 | 14.5 |
| ## 36 | 21.052043 | 18.9 |
| ## 41 | 35.110247 | 34.9 |
| ## 45 | 22.052093 | 21.2 |
| ## 51 | 20.047930 | 19.7 |
| ## 54 | 20.770800 | 23.4 |
| ## 74 | 23.976357 | 23.4 |
| ## 76 | 23.067717 | 21.4 |
| ## 78 | 21.720560 | 20.8 |
| ## 79 | 21.093743 | 21.2 |
| ## 82 | 24.860457 | 23.9 |
| ## 86 | 26.217757 | 26.6 |
| ## 92 | 22.767870 | 22.0 |
| ## 105 | 20.020050 | 20.1 |
| ## 108 | 19.483993 | 20.4 |
| ## 109 | 20.304587 | 19.8 |
| ## 111 | 20.827193 | 21.7 |
| ## 120 | 20.121900 | 19.3 |
| ## 127 | 16.186260 | 15.7 |
| ## 130 | 16.147150 | 14.3 |
| ## 131 | 20.730813 | 19.2 |
| ## 138 | 18.281033 | 17.1 |
| ## 142 | 13.493637 | 14.4 |
| ## 146 | 15.500093 | 13.8 |
| ## 151 | 20.015437 | 21.5 |
| ## 152 | 19.761130 | 19.6 |
| ## 155 | 18.222850 | 17.0 |
| ## 163 | 46.833493 | 50.0 |
| ## 167 | 47.814030 | 50.0 |
| ## 168 | 20.219743 | 23.8 |
| ## 170 | 22.561063 | 22.3 |
| ## 172 | 20.905493 | 19.1 |
| ## 178 | 24.074203 | 24.6 |
| ## 182 | 24.570117 | 36.2 |
| ## 184 | 27.017447 | 32.5 |
| ## 188 | 27.200040 | 32.0 |
| ## 198 | 33.071363 | 30.3 |
| ## 203 | 44.927280 | 42.3 |
| ## 205 | 47.201160 | 50.0 |
| ## 215 | 18.371497 | 23.7 |

| | | | |
|----|-----|-----------|------|
| ## | 218 | 23.833997 | 28.7 |
| ## | 221 | 28.319730 | 26.7 |
| ## | 224 | 25.110910 | 30.1 |
| ## | 244 | 24.638027 | 23.7 |
| ## | 246 | 19.322300 | 18.5 |
| ## | 247 | 21.748340 | 24.3 |
| ## | 250 | 25.842227 | 26.2 |
| ## | 252 | 27.598193 | 24.8 |
| ## | 255 | 22.529070 | 21.9 |
| ## | 257 | 41.468927 | 44.0 |
| ## | 262 | 41.117620 | 43.1 |
| ## | 271 | 21.095560 | 21.1 |
| ## | 293 | 28.219563 | 27.9 |
| ## | 294 | 22.710097 | 23.9 |
| ## | 300 | 33.396887 | 29.0 |
| ## | 305 | 33.659600 | 36.1 |
| ## | 307 | 33.720593 | 33.4 |
| ## | 312 | 23.620690 | 22.1 |
| ## | 316 | 20.005043 | 16.2 |
| ## | 320 | 21.677367 | 21.0 |
| ## | 323 | 21.595203 | 20.4 |
| ## | 326 | 24.908967 | 24.6 |
| ## | 330 | 23.183227 | 22.6 |
| ## | 348 | 24.221697 | 23.1 |
| ## | 352 | 25.583637 | 24.1 |
| ## | 355 | 20.084923 | 18.2 |
| ## | 357 | 15.484570 | 17.8 |
| ## | 370 | 40.163110 | 50.0 |
| ## | 378 | 13.501350 | 13.3 |
| ## | 393 | 11.563060 | 9.7 |
| ## | 394 | 14.905843 | 13.8 |
| ## | 401 | 9.587693 | 5.6 |
| ## | 403 | 11.907750 | 12.1 |
| ## | 405 | 9.486213 | 8.5 |
| ## | 406 | 10.370757 | 5.0 |
| ## | 410 | 14.278347 | 27.5 |
| ## | 411 | 24.342193 | 15.0 |
| ## | 417 | 11.015587 | 7.5 |
| ## | 422 | 15.316153 | 14.2 |
| ## | 445 | 11.282477 | 10.8 |
| ## | 449 | 14.532750 | 14.1 |
| ## | 453 | 17.220740 | 16.1 |
| ## | 455 | 13.849620 | 14.9 |
| ## | 472 | 20.849917 | 19.6 |
| ## | 481 | 21.442207 | 23.0 |
| ## | 484 | 20.732680 | 21.8 |
| ## | 486 | 21.714020 | 21.2 |
| ## | 487 | 19.572517 | 19.1 |
| ## | 490 | 13.506037 | 7.0 |
| ## | 493 | 19.939503 | 20.1 |
| ## | 495 | 20.397557 | 24.5 |
| ## | 496 | 19.780513 | 23.1 |